Assignment 2 Rubric (Postgraduate, 40 pts)

Summary

Component	Pts
A. Performance	**12**
B. Reproducibility and live verification	**6**
C. Method design and description	**8**
D. Controlled experiments, case study, and LLM judge validation	**8**
E. Report quality	**3**
F. Code quality	**3**
Total	**40**

Common Rules

Required interfaces. Internal architecture is flexible. Staff assess your code as submitted through the required commands and final input and output formats. Partial work is assessed under the relevant component bands.

Reported evidence. Reported results must be reproducible using the submitted code and configs, and match saved outputs. Normal variation from fresh LLM calls is allowed. A significant contradiction may result in zero for the affected reporting component. An isolated transcription error only removes the unsupported claim.

No hard-coding. Development-specific answers, support lists, option positions, or lookup tables are invalid methods. A result produced by these shortcuts receives zero for the affected performance track, regardless of its numerical score.

A. Performance (12 pts)

A.1 Development split (4 pts)

Each track contributes up to 2 points. Award the highest band reached.

Metric	1 pt	2 pts
Distribution score	$S \geq 0.30$	$S \geq 0.68$
MAP@100	MAP ≥ 0.20	MAP ≥ 0.44

Award zero for a score below the first cutoff, no valid submission, or hard-coded results.

A.2 Hidden-test split (8 pts)

Hidden-test marks use these fixed thresholds. Values on a threshold receive the higher band.

Distribution band	Pts	MAP@100 band	Pts
$S \geq 0.58$	5	MAP ≥ 0.61	3
$0.55 \leq S < 0.58$	4	$0.42 \leq$ MAP < 0.61	2
$0.50 \leq S < 0.55$	3	$0.20 \leq$ MAP < 0.42	1
$0.30 \leq S < 0.50$	2	MAP < 0.20, no valid submission, or hard-coded results	0
$0.26 \leq S < 0.30$	1		
$S < 0.26$, no valid submission, or hard-coded results	0		

B. Reproducibility and Live Verification (6 pts)

B.1 Automated mini reproducibility check (2 pts)

Pts	Descriptor
2	The submitted mini version of Best completes the indexing, retrieval, and judging workflow within three minutes and produces a valid final CSV.
1	At least one stage of your implementation produces usable output, but the complete workflow, time limit, or final CSV requirements are not satisfied.
0	Only unchanged B0 behaviour, or no stage of your implementation produces usable output.

B.2 Live demonstration (4 pts)

Only the first three minutes of `DEMO.mp4` are assessed.

Pts	Descriptor
4	Runs the submitted system and traces one question through retrieval, judging, and the final CSV. Gives a thorough, accurate explanation of how the relevant code and intermediate outputs produce the displayed result.
3	Shows a working run and gives a clear, substantially complete explanation of relevant code and outputs. Key connections are correct, but parts of the cross-stage trace or explanation of how the method works lack depth.
2	Shows relevant code and output with a broadly understandable explanation. Mainly describes the steps, with limited explanation of how the result was produced.
1	The demonstration is fragmented, substantially incomplete, or contains major explanatory errors. Connections between code and results are unclear, or the discussed result is not shown.
0	Missing, not the submitted work, edited rather than a continuous live take, or no meaningful live verification.

C. Method Design and Description (8 pts)

Data access and data analysis are each marked out of 4. Only the advanced methods in Main Specification Section 2.2 qualify. Each claimed method needs verifiable implementation code.

Pts	Descriptor for each area
4	A precise, code-consistent explanation of Best and at least one advanced method implemented during development, whether retained or rejected. Mechanisms, assumptions and the connections between method inputs and outputs within the pipeline are explained in depth. Key design choices and the reasons for retaining or rejecting each claimed advanced method are convincingly justified for the task.
3	Correct, specific explanation of Best and at least one implemented advanced method, whether retained or rejected. Main mechanisms and pipeline roles are clear and consistent with the code. Reasons for key choices are sound, but their justification for the task, including retained/rejected decisions, is only partly developed.
2	Substantially correct explanation of Best's core method and pipeline role. Choices are weakly justified, **or no qualifying advanced method is demonstrated**.

Pts	Descriptor for each area
1	Limited explanation of an implemented method, with some correct content that can be checked against the code, but important gaps or errors.
0	No verifiable core implementation, or the description is missing or contains no usable, correct explanation of the implemented method.

D. Controlled Experiments, Case Study, and LLM Judge Validation (8 pts)

D.1 Controlled experiments (4 pts)

Metric interpretation. For distribution, relate changes in predicted probabilities and answer-option ordering to the score, using the intuition in Appendix Section 4.1. Equations are not required. For MAP@100, explain how relevant-document ranks affect average precision. Analyse the metrics affected by the comparison.

Pts	Descriptor
4	All three comparisons isolate one method, component or setting change while keeping other factors fixed, and are reproducible. Each analysis convincingly links the isolated change to specific outputs and metric or efficiency results. The final method and configuration choices are well supported by the evidence, relevant trade-offs, and the effect of an observed method failure or limitation.
3	All three comparisons isolate one method, component or setting change while keeping other factors fixed and have verifiable run evidence. The explanations are clear and correct, but links between specific outputs, metric or efficiency principles, and the final choices are partly developed.
2	Two comparisons are developed with verifiable evidence, or all three have verifiable results but one or more do not isolate the selected change. The explanation is broadly correct, but gaps in controls or supporting outputs limit the conclusions.
1	At least one relevant comparison has verifiable run evidence, but the comparison design is weak or the analysis is mainly descriptive or substantially incorrect. The effect of the selected change is not clearly established.
0	No submitted run evidence, unrelated comparisons, or results that cannot be verified in important respects.

D.2 Case study (2 pts)

Pts	Descriptor
2	A convincing explanation of how ranked supports and two or three identified excerpts support or contradict the predicted probabilities. A limitation in the evidence or judging process is clearly linked to its possible effect on the prediction, with a well-justified, feasible improvement.
1	Uses relevant evidence, but the interpretation is largely descriptive, the limitation's effect is unclear, or the proposed improvement is weakly justified.
0	Missing, unrelated, or unsupported by submitted evidence.

D.3 LLM judge validation (2 pts)

Pts	Descriptor
2	Convincingly justifies the model choice by comparing two approved course-gateway LLMs using the same judging method and five or more manually labelled development examples. Labels and accuracy calculations are correct. Explains how errors, differences or agreement between the models support the choice.
1	Gives a relevant, sound reason for the model choice, but the accuracy comparison or its supporting evidence is incomplete or absent.
0	Missing, unrelated, or cannot be verified in important respects.

E. Report Quality (3 pts)

Only the first three pages of the report body are assessed. The cover page and Evidence Summary Table do not count towards this limit.

Pts	Descriptor
3	Clear, concise, and logically organised throughout. The opening overview clearly summarises Best's outputs and retained methods. Tables and text consistently explain the comparisons and final choice. The Evidence Summary Table accurately locates all assessed claims and results.
2	Generally clear and organised, with minor gaps in the overview, repetition, presentation issues, or weak connections between text and tables. The Evidence Summary Table is substantially complete and accurate. Reports exceeding the three-page body limit can score no higher than this band.
1	Substantial problems with organisation, explanation, formatting or tables make the report difficult to follow, or the Evidence Summary Table is present but substantially incomplete or inaccurate.
0	The Evidence Summary Table or report is missing, or the report cannot be meaningfully read or verified.

F. Code Quality (3 pts)

Only code you add or modify is assessed. Unchanged starter code is neither credited nor penalised. Integration problems caused by your changes are assessed.

Pts	Descriptor
3	Clear, maintainable code with consistent naming, appropriate reuse, and useful documentation of non-obvious logic. Configured choices work as described. Handling of relevant invalid inputs and model replies is appropriate, and integration with the supplied workflow is reliable.
2	Code is readable and can be assessed, with minor maintainability, documentation or defensive-programming issues.
1	Meaningful code is present, but unreliable integration, poor structure, extensive hard-coding, or very limited documentation makes it difficult to understand or verify.
0	Only unchanged starter behaviour, unsafe handling of API keys or other credentials, or code that cannot be meaningfully inspected or graded.